

**Лабораторная работа №6**  
**по курсу «Методы машинного обучения»**  
**на тему**  
**«Ансамбли моделей машинного обучения.»**

**Выполнил:**  
**Хотин П.Ю.**  
**ИУ5-24М**

Москва, 2020 год

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor
from sklearn.metrics import mean_absolute_error, accuracy_score,
r2_score

data = pd.read_csv("advertising.csv")

data.head()

```

	TV	Radio	Newspaper	Sales
1	230.1	37.8	69.2	22.1
2	44.5	39.3	45.1	10.4
3	17.2	45.9	69.3	9.3
4	151.5	41.3	58.5	18.5
5	180.8	10.8	58.4	12.9

```

data_X = data[["TV", "Radio", "Newspaper"]]

data_X

```

	TV	Radio	Newspaper
1	230.1	37.8	69.2
2	44.5	39.3	45.1
3	17.2	45.9	69.3
4	151.5	41.3	58.5
5	180.8	10.8	58.4
..	...	...	...
196	38.2	3.7	13.8
197	94.2	4.9	8.1
198	177.0	9.3	6.4
199	283.6	42.0	66.2
200	232.1	8.6	8.7

200 rows  $\times$  3 columns

```
data_Y = data[["Sales"]]
```


```
X_train, X_test, y_train, y_test = train_test_split(
    data_X, data_Y, test_size=0.25, random_state=1)
```

*# Качество отдельных моделей*

```
def val_mae(model):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    plt.plot(X_test, y_test, 'g.')
    plt.plot(X_test, y_pred, 'ro')
    plt.show()
    result = mean_absolute_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    print(model)
    print('MAE={}'.format(result))
    print('R2={}'.format(r2))
```

```
for model in [
    GradientBoostingRegressor(),
    RandomForestRegressor(n_estimators=50)
]:
    val_mae(model)
    print('=====\n\n')
```


```
/usr/local/lib/python3.7/site-packages/sklearn/ensemble/_gb.py:1454:
DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
```



```
/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:3:
DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples,), for
example using ravel().
```

This is separate from the ipykernel package so we can avoid doing imports until

```
GradientBoostingRegressor(alpha=0.9, ccp_alpha=0.0,
criterion='friedman_mse',
                                init=None, learning_rate=0.1, loss='ls',
max_depth=3,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0,
min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0,
n_estimators=100,
                                n_iter_no_change=None, presort='deprecated',
                                random_state=None, subsample=1.0,
tol=0.0001,
                                validation_fraction=0.1, verbose=0,
warm_start=False)
MAE=0.48997309191670874
R2=0.9831579266623767
=====
```



```

RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                      max_depth=None, max_features='auto',
                      max_leaf_nodes=None,
                      max_samples=None, min_impurity_decrease=0.0,
                      min_impurity_split=None, min_samples_leaf=1,
                      min_samples_split=2,
                      min_weight_fraction_leaf=0.0,
                      n_estimators=50, n_jobs=None, oob_score=False,
                      random_state=None, verbose=0, warm_start=False)
MAE=0.5131199999999992
R2=0.9815468551914713
=====

```

## Модель градиентного бустинга показала лучший результат на тестовой выборке

```

from sklearn.model_selection import RandomizedSearchCV

n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000,
num = 10)]

max_features = ['auto', 'sqrt']

max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
max_depth.append(None)

min_samples_split = [2, 5, 10]

min_samples_leaf = [1, 2, 4]

bootstrap = [True, False]

random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,

```



```

criterion='mse',
max_depth=None,

max_features='auto',

max_leaf_nodes=None,

max_samples=None,

min_impurity_decrease=0.0,

min_impurity_split=None,

min_samples_leaf=1,

min_samples_split=2,

min_weight_fraction_leaf=0.0,

n_estimators=100,
n_jobs=None,
oob_score=False...
    param_distributions={'bootstrap': [True, False],
                          'max_depth': [10, 20, 30, 40,
50, 60,
70, 80, 90, 100,
110,
None],
                          'max_features': ['auto',
'sqrt'],
                          'min_samples_leaf': [1, 2, 4],
                          'min_samples_split': [2, 5,
10],
                          'n_estimators': [200, 400,
600, 800,
1000, 1200,
1400, 1600,
1800, 2000]}},
    pre_dispatch='2*n_jobs', random_state=42,
refit=True,
    return_train_score=False, scoring=None, verbose=2)
rf_random.best_params_

```

```
{'n_estimators': 800,
  'min_samples_split': 2,
  'min_samples_leaf': 1,
  'max_features': 'auto',
  'max_depth': 100,
  'bootstrap': True}
```

```
def evaluate(model, test_features, test_labels):
    predictions = model.predict(test_features)
    error = mean_absolute_error(y_test, predictions)
    r2 = r2_score(y_test, predictions)
    print('Model Performance')
    print('MAE: {:.4f}'.format(error))
    print('R2 score: {:.4f}'.format(r2))
    print('=====\n\n')
```

```
base_model = RandomForestRegressor(n_estimators = 10, random_state =
42)
```

```
base_model.fit(X_train, y_train)
evaluate(base_model, X_test, y_test)
```

Model Performance

MAE: 0.5994

R2 score: 0.9713

=====

```
/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:11:
DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples,), for
example using ravel().
```

```
# This is added back by InteractiveShellApp.init_path()
```

```
best_random = rf_random.best_estimator_
evaluate(best_random, X_test, y_test)
```

Model Performance

MAE: 0.5178



R2 score: 0.9820

=====

Видно, что подбор гиперпараметров улучшил нашу модель, уменьшив ошибку на 0.08

```
n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000,
num = 10)]
```

```
max_features = ['auto', 'sqrt']
```

```
max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
max_depth.append(None)
```

```
min_samples_split = [2, 5, 10]
```

```
min_samples_leaf = [1, 2, 4]
```

```
bootstrap = [True, False]
```

```
random_grid_Booster = {'n_estimators': n_estimators,
                        'max_features': max_features,
                        'max_depth': max_depth,
                        'min_samples_split': min_samples_split,
                        'min_samples_leaf': min_samples_leaf,
                        }
```

```
gb = GradientBoostingRegressor()
```

```
gb_random = RandomizedSearchCV(estimator = gb, param_distributions =
random_grid_Booster, n_iter = 100, cv = 3, verbose=2, random_state=42,
n_jobs = -1)
```

```
gb_random.fit(X_train, y_train)
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 16 concurrent
workers.
```

Fitting 3 folds for each of 100 candidates, totalling 300 fits

```
[Parallel(n_jobs=-1)]: Done 9 tasks      | elapsed: 0.4s
[Parallel(n_jobs=-1)]: Done 221 tasks   | elapsed: 7.7s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed: 10.2s finished
/usr/local/lib/python3.7/site-packages/sklearn/ensemble/_gb.py:1454:
DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
```

```
RandomizedSearchCV(cv=3, error_score=nan,
                   estimator=GradientBoostingRegressor(alpha=0.9,
ccp_alpha=0.0,

criterion='friedman_mse',

init=None,

learning_rate=0.1,

loss='ls',

max_depth=3,

max_features=None,

max_leaf_nodes=None,

min_impurity_decrease=0.0,

min_impurity_split=None,

min_samples_leaf=1,

min_samples_split=2,

min_weight_fraction_leaf=0.0,
```

```

n_estimators=100,
                                n_...
                                iid='deprecated', n_iter=100, n_jobs=-1,
                                param_distributions={'max_depth': [10, 20, 30, 40,
50, 60,
                                70, 80, 90, 100,
110,
                                None],
                                'max_features': ['auto',
'sqrt'],
                                'min_samples_leaf': [1, 2, 4],
                                'min_samples_split': [2, 5,
10],
                                'n_estimators': [200, 400,
600, 800,
                                1000, 1200,
1400, 1600,
                                1800, 2000]}},
                                pre_dispatch='2*n_jobs', random_state=42,
refit=True,
                                return_train_score=False, scoring=None, verbose=2)

gb_random.best_params_
{'n_estimators': 1400,
 'min_samples_split': 10,
 'min_samples_leaf': 2,
 'max_features': 'auto',
 'max_depth': 40}

def evaluate(model, test_features, test_labels):
    predictions = model.predict(test_features)
    error = mean_absolute_error(y_test, predictions)
    r2 = r2_score(y_test, predictions)
    print('Model Performance')
    print('MAE: {:.4f}'.format(error))
    print('R2 score: {:.4f}'.format(r2))
    print('=====\n\n')

base_model = GradientBoostingRegressor()

```

```
base_model.fit(X_train, y_train)
evaluate(base_model, X_test, y_test)
```

Model Performance

MAE: 0.4890

R2 score: 0.9832

=====

```
/usr/local/lib/python3.7/site-packages/sklearn/ensemble/_gb.py:1454:
DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
```

```
    y = column_or_1d(y, warn=True)
```

```
best_random = gb_random.best_estimator_
evaluate(best_random, X_test, y_test)
```

Model Performance

MAE: 0.5078

R2 score: 0.9839

=====

**Подбор параметров в градиентном бустинге не дал прироста качества (оно и так в целом было достаточно высокое)**